

POINTY

The AI that shows you where to click

Pointy is an always-on desktop assistant that watches your screen only when you ask it to, answers what you're stuck on out loud, and points — literally, with an animated on-screen dot — at exactly what to click. A second, always-free feature reports how you spend time on your own machine, no cloud AI required.

Tauri + Rust

Local-first

Voice-native

This brief covers the concept, the full wake-word-to-voice pipeline, the Tauri tech stack, realistic API cost modeling, the competitive landscape, a three-tier pricing model, go-to-market, an MVP roadmap, key risks, and the privacy approach — everything needed to take Pointy from idea to a priced, buildable product.

Contents

- 1 Executive Summary**
Why Pointy, and why now.
- 2 Core Concept**
The two features, in plain terms.
- 3 Architecture & Pipeline**
Wake word → STT → vision LLM → voice + overlay.
- 4 Tech Stack**
Every layer, the choice, and the reasoning.
- 5 API Cost Model**
Real per-user cost at light, moderate, and heavy usage.
- 6 Competitive Landscape**
Cluely and the rest of the AI-overlay category.
- 7 Positioning & Differentiation**
Where Pointy should win.
- 8 Pricing Model**
Three tiers, sized against real cost.
- 9 Revenue Projection**
A simple early-stage scenario.
- 10 Go-to-Market**
Where the first 1,000 users come from.
- 11 MVP Roadmap**
What to build first, second, third.
- 12 Risks & Mitigations**
What could go wrong, and the plan for it.
- 13 Privacy & Security**
The approach and the promise to users.

1 Executive Summary

Voice assistants can talk. Chatbots can see a screenshot. Almost nothing combines the two into a tool that watches what you're doing, understands where you're stuck, and shows you — not just tells you — what to do next. That gap is Pointy's opening.

The nearest existing product, Cluely, proves the category has real demand — people already pay \$20–150/month for an AI that watches their screen and talks back. But Cluely is built for a fundamentally different job: hiding AI assistance from another person on a call. Pointy is built for the opposite job: helping you, alone, get unstuck in your own software, visibly and transparently.

The core bet

A local-first, voice-native "show me where to click" assistant can undercut Cluely on trust and price while building a genuinely harder, more defensible feature: precise on-screen guidance, not just an answer in a text box.

2 Core Concept

Pointy runs quietly in the background, woken by a hotkey or a spoken phrase ("hey Pointy"). Once active, it listens for a spoken question, reads the current screen, and responds by voice while animating a dot over the exact UI element it's referring to — a settings menu, an icon, a button buried three panels deep.

Two features, one always free

- **Screen help** — "How do I change the icon pack in VS Code?" → Pointy reads the screen, answers out loud, and animates a dot guiding you to the setting.
- **Usage awareness** — "How long did I spend on YouTube?" → answered instantly from local activity tracking, no API call needed. Included on every tier, including free.

3 Architecture & Pipeline

Four steps, only one of which costs API money per use. The other three run locally, keeping baseline cost near zero and giving Pointy a real privacy story.

Step	What it does	Where it runs
1. Wake word listener	Listens for the activation phrase or hotkey	Local — near-zero cost
2. Speech-to-text	Transcribes the spoken question	Local (or cheap cloud STT)
3. Screen capture + vision LLM	Screenshot + question sent to a vision-capable model for understanding	Cloud API — the paid step
4. Voice answer + overlay	Speaks the answer; draws the guide-dot using OS accessibility APIs for precise coordinates	Local

Key implementation challenge

The vision model gives an approximate description of where something is, not pixel-perfect coordinates. Pointy pairs the model's answer with the OS accessibility layer (UI Automation on Windows, the Accessibility API on macOS) to find the real bounding box of the element and place the dot precisely — this is the hardest engineering problem in the product, and the biggest differentiator from a simple “AI chatbot with a screenshot.”

4 Tech Stack

Layer	Choice	Why
App shell	Tauri (Rust + webview)	Small binary, low background overhead, direct OS API access for overlays, hotkeys, screen capture
Wake word	Porcupine (Picovoice) or openWakeWord	Runs fully on-device, negligible CPU, no audio leaves the machine until triggered
Speech-to-text	whisper.cpp (local) or Deepgram/AssemblyAI (cloud)	Local keeps marginal cost at zero; cloud trades a few cents/hour for better accuracy
Screen understanding	Claude vision API — Haiku by default, Sonnet on escalation	Model routing keeps the one paid step cheap without sacrificing quality on hard questions

Layer	Choice	Why
Text-to-speech	Piper (local) or a premium cloud voice on the top tier	Local keeps free/entry tiers sustainable; premium voice is a real upsell
Overlay + highlighting	Transparent always-on-top Tauri window + OS accessibility APIs	Precise element targeting instead of approximate model coordinates
Usage tracking	Local window/tab focus logging + local DB (SQLite)	Zero API cost; works fully offline; only calls an LLM for open-ended questions about the data

5 API Cost Model

Only the screen-understanding step carries meaningful per-query cost. Using Claude's per-token pricing, a single query — one screenshot plus a short question in, a short spoken answer out — costs roughly \$0.003–\$0.01 on Haiku-tier vision, or \$0.01–\$0.02 on Sonnet-tier for harder questions.

Usage level	Model mix	Est. cost / user / month
Light — ~8 queries/day	Haiku only	\$0.30 – \$0.50
Moderate — ~40 queries/day	Haiku, escalate to Sonnet on hard cases	\$3 – \$5
Heavy — ~100–150 queries/day	Sonnet default + premium voice	\$8 – \$12

Margin lever

Model routing — Haiku by default, Sonnet only on ambiguous or visually dense screens — is the single biggest cost control available. It separates a \$3/month user from a \$12/month user at the same query volume, and it's a software decision, not an infrastructure investment.

6 Competitive Landscape

Cluely is the closest existing product: a desktop overlay that reads your screen via OCR, captures audio, and surfaces AI answers over live calls, with a GPU-level invisible overlay as its flagship feature. The wider category also includes meeting-note tools and general AI assistants that touch some of the same ground without directly competing on Pointy's core job.

Product	What it actually does	Overlap with Pointy
Cluely	Live-call AI answers, hidden from the other party via a GPU-level invisible overlay	High — closest architecture, opposite use case
Rewind.ai / similar recall tools	Continuously records screen + audio, lets you search your own past activity later	Low — retrospective search, not live in-the-moment guidance
Otter.ai / Granola	Meeting transcription and note-taking during calls	Low — focused on notes, not screen navigation

Product	What it actually does	Overlap with Pointy
General AI chat apps	Answer a pasted screenshot or question on request, no always-on listening or overlay targeting	Medium — same underlying model capability, none of the always-on or pointing UX

7 Positioning & Differentiation

	Cluely	Pointy
Core use case	Live-call answers, often for interviews/sales, hidden from the other party	Personal software navigation help, visible only to you by design
Signature tech	GPU-level overlay invisible to screen-share software	Precise UI-element targeting via OS accessibility APIs + animated guide-dot
Processing	Cloud-dependent OCR, STT, and LLM	Local-first: wake word, STT, and TTS run on-device; only screen understanding hits the cloud
Trust track record	2025 breach exposed 83,000 users' data; documented billing/refund complaints	Local-first processing and transparent flat pricing as the core pitch, not an afterthought
Pricing top tier	\$75–150/month, gated on the invisibility feature	\$25/month, gated on usage volume and model quality, not on hiding something

Where Pointy should win

- **Different trust story** — a personal accessibility tool has none of the ethical baggage of a call-deception tool; market it as help for when you're stuck, not stealth.
- **Security & billing as the pitch** — local-first processing and clear flat pricing directly answer Cluely's most public failures.
- **Skip the GPU-hook overlay entirely** — it's expensive engineering built for a hiding use case Pointy doesn't have.
- **Depth over breadth** — actually pointing at the right element, not just describing it, is a defensible wedge feature competitors haven't built.

8 Pricing Model

Three tiers, sized against the cost curve above to hold 55–75% gross margin, plus a bring-your-own-key option for early, technical adopters.

	Free	Pro — \$12/mo	Power — \$25/mo
Screen queries	~8/day	~40/day soft cap	Fair-use unlimited (~150/day)
Model	Haiku-tier vision	Haiku, escalates to Sonnet	Sonnet by default
Voice	OS built-in TTS	Local neural TTS	Premium cloud voice
Usage tracking	Included, unlimited	+ natural-language Q&A;	+ weekly digest & trend alerts
Your est. cost/user	~\$0.30–\$0.50/mo	~\$3–\$5/mo (~65–75% margin)	~\$8–\$12/mo (~55–65% margin)

Additional pricing notes

- **Bring-your-own-key tier — \$5/mo:** for technical early adopters. They supply their own Anthropic API key and pay Anthropic directly; you charge a flat fee for the app itself with zero API cost exposure. Good for early traction before usage-metering infrastructure exists.
- **Annual discount:** ~20% off (e.g., Pro at \$120/yr instead of \$144) — standard SaaS practice, improves cash flow and cuts churn.
- **Quiet fair-use caps on “unlimited”:** protects margin from outlier heavy users without needing to advertise a hard limit, avoiding the billing-dispute pattern seen with competitors.
- **Usage tracking is free on every tier, including Free:** it costs nothing to run locally and drives daily habit formation ahead of conversion to a paid tier.

9 Revenue Projection

A simple early-stage scenario — not a forecast, a sanity check on whether the tier structure produces a sustainable business at modest scale. Assumes a typical freemium conversion pattern: most users stay free, a minority convert to Pro, a small minority convert to Power.

Tier	Users	Price	Revenue	Est. API cost
Free	3,500	\$0	\$0	~\$1,400
Pro	380	\$12/mo	\$4,560	~\$1,520
Power	70	\$25/mo	\$1,750	~\$700
Total (MRR)	3,950	—	\$6,310	~\$3,620

Reading this scenario

At roughly 4,000 total users with a ~11% paid conversion rate (typical for freemium productivity tools), gross margin on the paid base alone is around 55–65%, before accounting for the free tier's cost as a customer-acquisition expense. The main levers to improve this: raise conversion through better onboarding, or shift the free tier's query cap down slightly once usage data shows where casual users actually land.

10 Go-to-Market

Pointy's first users are almost certainly developers and power users — people already comfortable installing a background utility, and the exact audience who'll appreciate a VS Code-navigation demo. Early channels should lean into that.

- **Launch on Product Hunt and Hacker News** with a demo video centered on the VS Code icon-pack example — concrete, relatable “I was stuck, it pointed at the answer” moments convert far better than abstract pitches.
- **Short demo videos on YouTube/X/TikTok** showing the dot animation in action — this is a visual product, and the guide-dot is the single most shareable five seconds of the experience.
- **Developer communities** (Reddit's r/programming, dev Discord servers, Indie Hackers) for direct feedback and the bring-your-own-key tier, which removes the trust barrier for a first install.
- **Accessibility and productivity communities** as a secondary wedge — the same “show me where to click” mechanic is genuinely useful for users with cognitive load or motor-control difficulty navigating complex software, a market Cluely doesn't touch at all.

11 MVP Roadmap

Phase	Scope	Goal
Phase 1 — MVP	Hotkey activation, screenshot + question to Haiku, spoken answer. No dot overlay yet.	Validate that voice Q&A; about the screen is useful enough to keep using daily
Phase 2 — Guide-dot	Add OS accessibility integration and the animated overlay pointing at real UI elements	Ship the actual differentiator; this is what makes Pointy, Pointy
Phase 3 — Wake word + usage tracking	Always-on local wake word listener; local window/tab tracking and the usage-question feature	Make the product ambient rather than manually triggered; add a free daily-habit hook
Phase 4 — Tiers + billing	Model routing (Haiku/Sonnet), subscription billing, fair-use caps	Turn on monetization once usage patterns from Phases 1–3 are known
Phase 5 — Expansion	Multi-monitor support, per-app custom skills (IDEs, design tools), team/enterprise tier	Grow beyond individual users once the core loop is proven

12 Risks & Mitigations

Risk	Why it matters	Mitigation
Accessibility API fragmentation	Windows UI Automation and macOS Accessibility behave differently per-app; some apps expose poor accessibility trees	Start with a curated list of well-supported apps (VS Code, browsers, Office); fall back to an approximate dot position when accessibility data is missing
Screen-recording permission friction	macOS and Windows both require explicit, sometimes repeated, user consent for screen capture	Design the first-run flow around this explicitly; treat it as an onboarding step, not an error state
Heavy-user cost outliers	A small number of power users on flat-rate plans can erode margin	Quiet fair-use caps and model routing, reviewed monthly against real usage data
Competitive response	An incumbent (Cluely or a new entrant) could copy the guide-dot mechanic	Move fast on the accessibility-API integration — it's the hard part to replicate well, not the idea itself

Risk	Why it matters	Mitigation
Trust and data handling	The product reads screens by nature; a single mishandled-data incident would be existential, as it was for Cluely	Local-first architecture, minimal retention, and a clear published data policy from day one

13 Privacy & Security

Because Pointy reads screens by design, privacy and security aren't a compliance checkbox — they're the actual product pitch, and the most direct answer to Cluely's public failures.

- **Local-first by default** — wake word detection, transcription, and voice output all run on-device. Only the screenshot and question for the current query are sent to the cloud, and only when the user actively triggers a request.
- **No standing screen access** — Pointy captures a screenshot at the moment of a query, not a continuous recording. There is no rolling buffer of screen history to secure or leak.
- **Clear, published retention policy** — screenshots and queries are not stored beyond what's needed to generate the answer, and this should be stated plainly in-product, not buried in a terms-of-service document.
- **No credential storage in plaintext, no public-repo exposure risk** — the exact failure mode that caused Cluely's 2025 breach should be an explicit item in the security checklist before launch, not an afterthought.
- **Transparent billing** — clear cancellation flow, no silent annual auto-upgrades, and a real refund policy — directly addressing the billing complaints that show up repeatedly in Cluely's user reviews.

The trust pitch, in one line

"Your screen doesn't need to leave your machine to help you — and when it does, for a moment, you'll know exactly why."

Prepared as a working product brief. Pricing and cost figures are estimates based on published API rates and competitor pricing as of August 2026, and should be revisited as real usage data comes in.